

EN.600.475 Machine Learning

Linear Regression

Raman Arora
Lecture 4
February 8, 2017

- Loss and Risk
- Least squares estimation

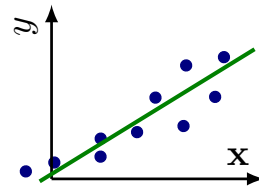
Slides credit: Greg Shakhnarovich

Review

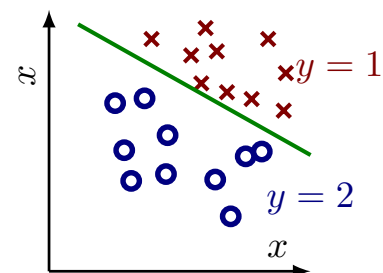
Review: supervised learning

- Task: build a mapping from input \mathcal{X} output \mathcal{Y}
- Given a training set (\mathbf{x}_i, y_i) , $i = 1, \dots, N$, with $\mathbf{x}_i \in \mathcal{X}$, $y_i \in \mathcal{Y}$.
- Goal (informally): predict y accurately for future x s

regression: $\mathcal{Y} = \mathbb{R}$
learn a (continuous) function f



classification: $\mathcal{Y} = \{1, \dots, C\}$
learn a separator between classes



Review: supervised learning pipeline

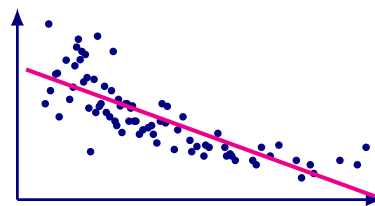
- **Set up (define)** a supervised learning problem
- **Data collection** for training and test set.
- **Representation** choose how data are fed to the model
- **Modeling** Choose a *hypothesis (model) class*
- **Estimation (learning)** Find best hypothesis you can in the chosen class, given the data.
- **Model (class) selection**

3

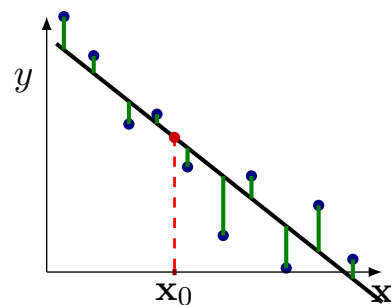


Review: linear regression

- Two goals in mind:
Explain the data
Make predictions



- Model class: linear functions
- Fitting criterion, to guide selection of a function: sum of squared distances from data to the line, along y axis



4



Roadmap

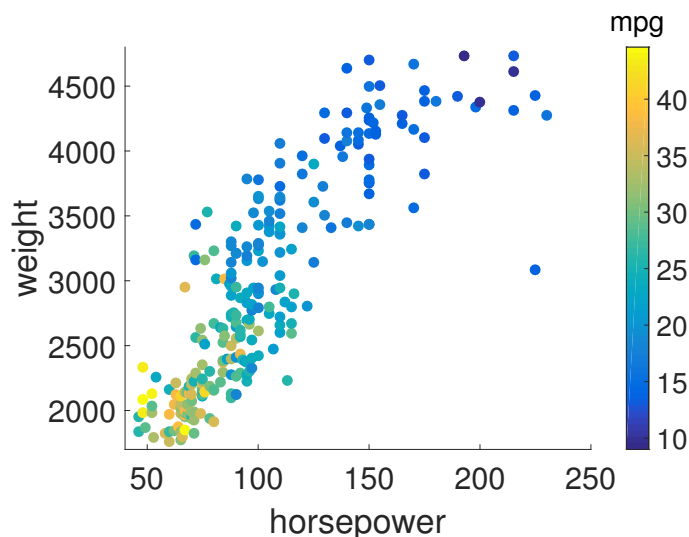
- General form of linear regression and least squares fit
- Loss and risk: definitions and analysis
- Analysis of error in empirical risk minimization

5



Multiple input variables

- Can consider additional features; e.g., x_1 horsepower and x_2 vehicle weight.
- We now have mapping from $\mathbf{x} = (x_1, x_2) \in \mathbb{R}^2$ to y



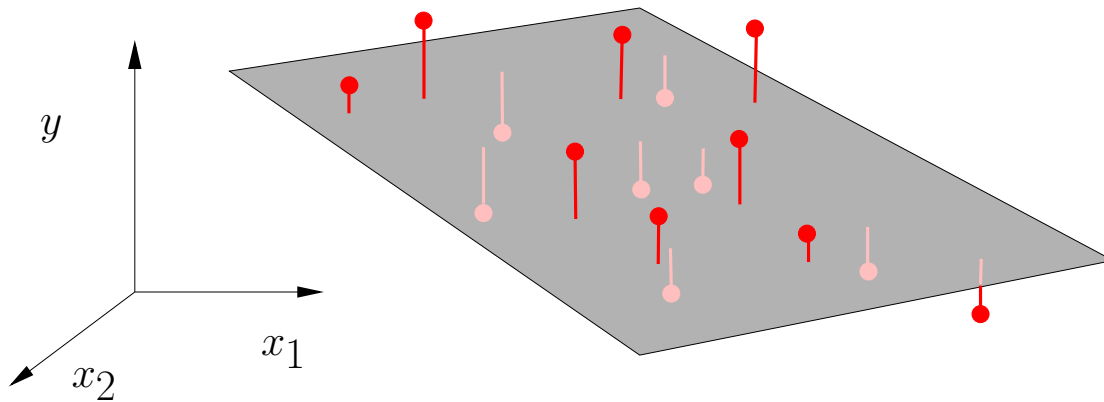
6

colorbar: one possible way to convey multi-dimensional plots



Fitting a plane to data

- Can use the same criterion: minimize sum of square distances along y -axis



7



Linear functions

- General form:

$$\begin{aligned} f(\mathbf{x}; \mathbf{w}) &= w_0 + w_1 x_1 + \dots + w_d x_d \\ &= \mathbf{w} \cdot \mathbf{x} \end{aligned}$$

denoting $x_0 \equiv 1$

- 1D case ($\mathcal{X} = \mathbb{R}$): a line
- $\mathcal{X} = \mathbb{R}^2$: a plane
- *Hyperplane* in general, d -D case.

8



Notation

We will mostly stick to these throughout the course:

- \mathbf{x}_i the i -th data point in \mathcal{X} (column vector)
Often $\mathcal{X} \equiv \mathbb{R}^d$, so that $\mathbf{x}_i = [x_{i1}, \dots, x_{id}]$
Often assume also $x_{i0} \equiv 1$
- y_i the label of the i -th data point; $y_i \in \mathcal{Y}$
- \mathbf{x}_0, y_0 a single test point and its (unknown) label
- \mathbf{X} the $N \times d$ data matrix where i -th row is \mathbf{x}_i
- \mathbf{y} the label vector $\mathbf{y} = [y_1, \dots, y_N]$
- $\mathbf{w} \cdot \mathbf{x}$ inner (dot) product, $\sum_j w_j x_j$
sometimes write $\mathbf{w}^T \mathbf{x}$

9



Loss function

- Recall: target labels are in \mathcal{Y} (e.g., regression: $\mathcal{Y} \equiv \mathbb{R}$)
- A *loss function* $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ maps prediction to cost, given true value:
 $\ell(\hat{y}, y)$ defines the penalty paid for predicting \hat{y} when the true value is y .
- Standard choice for regression: squared loss $\ell(\hat{y}, y) = (\hat{y} - y)^2$
is it a good loss function?..
- It is symmetric (sign of mistake doesn't matter); non-negative; gives zero loss for correct prediction
- Vaguely justifiable as “energy” of something
- Penalizes quite harshly for larger mistakes

10



Empirical loss

- We consider a *parametric* function $f(\mathbf{x}; \mathbf{w})$
E.g., linear function: $f(\mathbf{x}; \mathbf{w}) = w_0 + \sum_{j=1}^d w_j x_{ij} = \mathbf{w} \cdot \mathbf{x}_i$
- The *empirical loss* of function $y = f(\mathbf{x}; \mathbf{w})$ on a set \mathbf{X} :

$$L(\mathbf{w}, \mathbf{X}, \mathbf{y}) = \frac{1}{N} \sum_{i=1}^N \ell(f(\mathbf{x}_i; \mathbf{w}), y_i)$$

- LSQ minimizes the empirical loss when ℓ is squared loss.
- We care about accuracy of *predicting* labels for new examples.
Why/when does empirical loss minimization help us achieve that?

11



Loss: empirical and expected

- Fundamental assumption: example \mathbf{x} /label y are drawn from a joint probability distribution $p(\mathbf{x}, y)$.
- Data are i.i.d.: same (unknown!) distribution for all pairs (\mathbf{x}, y) in both training and test data.
- We can measure the empirical loss on training set

$$L(\mathbf{w}, \mathbf{X}, \mathbf{y}) = \frac{1}{N} \sum_{i=1}^N \ell(f(\mathbf{x}_i; \mathbf{w}), y_i)$$

- The ultimate goal is to minimize the *expected loss*, also known as *risk*:

$$R(\mathbf{w}) = E_{(\mathbf{x}_0, y_0) \sim p(\mathbf{x}, y)} [\ell(f(\mathbf{x}_0; \mathbf{w}), y_0)]$$

12



Loss: empirical and expected

- Empirical loss:

$$L(\mathbf{w}, \mathbf{X}, \mathbf{y}) = \frac{1}{N} \sum_{i=1}^N \ell(f(\mathbf{x}_i, \mathbf{w}), y_i)$$

- Risk:

$$R(\mathbf{w}) = E_{(\mathbf{x}_0, y_0) \sim p(\mathbf{x}, y)} [\ell(f(\mathbf{x}_0, \mathbf{w}), y_0)]$$

- Empirical risk minimization (ERM) approach: to the extent that the training set is a representative of the underlying distribution $p(\mathbf{x}, y)$, the empirical loss serves as a proxy for the risk (expected loss).
- Technically: estimate $p(\mathbf{x}, y)$ by the *empirical distribution* of data.

13



Learning via empirical loss minimization

Two steps:

- Select a restricted class \mathcal{H} of *hypotheses* $f : \mathcal{X} \rightarrow \mathcal{Y}$
Here: linear functions parametrized by \mathbf{w} : $f(\mathbf{x}; \mathbf{w}) = \mathbf{w} \cdot \mathbf{x}$
- Select a hypothesis $f^* \in \mathcal{H}$ based on training set (X, Y)
Here: minimize empirical squared loss, i.e., select $f(\mathbf{x}; \mathbf{w}^*)$ where

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

- How do we find $\mathbf{w}^* = [w_0^*, w_1^*, \dots, w_d^*]$?

14



Least squares: estimation

- We need to minimize L w.r.t. \mathbf{w}

$$L(\mathbf{w}, \mathbf{X}, \mathbf{y}) = L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

- Necessary condition to minimize L :

$$\frac{\partial L(\mathbf{w}, \mathbf{X}, \mathbf{y})}{\partial \mathbf{w}} = \mathbf{0},$$

i.e., derivatives w.r.t. w_0, w_1, \dots, w_d must all be zero.

15



Matrix derivatives

- Scalar valued function of one variable

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad \text{derivative: } \frac{df}{dx}$$

- Scalar valued function of multiple scalar variables

$$f : \underbrace{\mathbb{R} \times \dots \times \mathbb{R}}_{d \text{ times}} \rightarrow \mathbb{R} \quad \text{gradient: } \nabla f = \left[\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_d} \right]$$

- If we collect multiple variables in a vector: $\mathbf{x} \in \mathbb{R}^d$:

$$\nabla f = \frac{\partial f}{\partial \mathbf{x}}$$

derivative of f w.r.t. \mathbf{x} has the same dimension as \mathbf{x}

16



Least squares: estimation

$$L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

$$\frac{\partial}{\partial w_0} L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \frac{\partial}{\partial w_0} (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2 = 0$$

$$\Rightarrow \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) = 0$$

- $y_i - \mathbf{w} \cdot \mathbf{x}_i$ is the *prediction error* on the i -th example.
- \Rightarrow Necessary condition for optimal \mathbf{w} is that the errors have zero mean. (why does it make sense?)

17



Least squares: estimation

$$L(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

$$\frac{\partial}{\partial w_j} L(\mathbf{w}) = -\frac{2}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) x_{ij} = 0.$$

- Second necessary condition: errors are *uncorrelated* with the data!
(And with *any linear function* of the data)
- $d + 1$ linear equations in $d + 1$ unknowns w_0, w_1, \dots, w_d

$$\sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) x_{ij} = 0 \quad \forall j = 1, \dots, d, \quad (1)$$

$$\sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) = 0 \quad (2)$$

18



Least squares in matrix form

$$\mathbf{X} = \begin{bmatrix} 1 & x_{11} & \cdots & x_{1d} \\ \vdots & & & \\ 1 & x_{N1} & \cdots & x_{Nd} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} w_0 \\ \vdots \\ w_d \end{bmatrix}$$

- Predictions: $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$, errors: $\mathbf{y} - \mathbf{X}\mathbf{w}$, empirical loss:

$$\begin{aligned} L(\mathbf{w}, \mathbf{X}, \mathbf{y}) &= \frac{1}{N} (\mathbf{y} - \mathbf{X}\mathbf{w}) \cdot (\mathbf{y} - \mathbf{X}\mathbf{w}) = \frac{1}{N} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) \\ &= \frac{1}{N} (\mathbf{y}^T - \mathbf{w}^T \mathbf{X}^T) (\mathbf{y} - \mathbf{X}\mathbf{w}) \end{aligned}$$

Using $(AB)^T = B^T A^T$, $(A + B)^T = A^T + B^T$, $(A^T)^T = A$.

19



Derivative of loss

$$L(\mathbf{w}) = \frac{1}{N} (\mathbf{y}^T - \mathbf{w}^T \mathbf{X}^T) (\mathbf{y} - \mathbf{X}\mathbf{w}).$$

$$\frac{\partial \mathbf{a}^T \mathbf{b}}{\partial \mathbf{a}} = \frac{\partial \mathbf{b}^T \mathbf{a}}{\partial \mathbf{a}} = \mathbf{b}, \quad \frac{\partial \mathbf{a}^T \mathbf{B} \mathbf{a}}{\partial \mathbf{a}} = 2\mathbf{B}\mathbf{a}$$

$$\begin{aligned} \frac{\partial L(\mathbf{w})}{\partial \mathbf{w}} &= \frac{1}{N} \frac{\partial}{\partial \mathbf{w}} [\mathbf{y}^T \mathbf{y} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}] \\ &= \frac{1}{N} [\mathbf{0} - \mathbf{X}^T \mathbf{y} - (\mathbf{y}^T \mathbf{X})^T + 2\mathbf{X}^T \mathbf{X} \mathbf{w}] \\ &= -\frac{2}{N} (\mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X} \mathbf{w}) \end{aligned}$$

20



Least squares solution

$$\frac{\partial}{\partial \mathbf{w}} L(\mathbf{w}) = -\frac{2}{N} (\mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X} \mathbf{w}) = 0$$

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \mathbf{X} \mathbf{w} \Rightarrow \mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- $\mathbf{X}^\dagger \triangleq (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$ is called the *Moore-Penrose pseudoinverse* of \mathbf{X} .
- Linear regression in Python:

```
X[:,0]=1; X[:,1:]=x # assumes X is right size
w=np.dot(np.linalg.pinv(X),y)
```
- Prediction: $\hat{\mathbf{y}} = \mathbf{y}^T \mathbf{X}^\dagger \mathbf{x}_0$

Note: we have $d + 1$ numbers in \mathbf{w}^* capture what the training data \mathbf{X}, \mathbf{y} tell us about $\mathcal{X} \rightarrow \mathcal{Y}$ *under our model class*

21



Data set size and regression

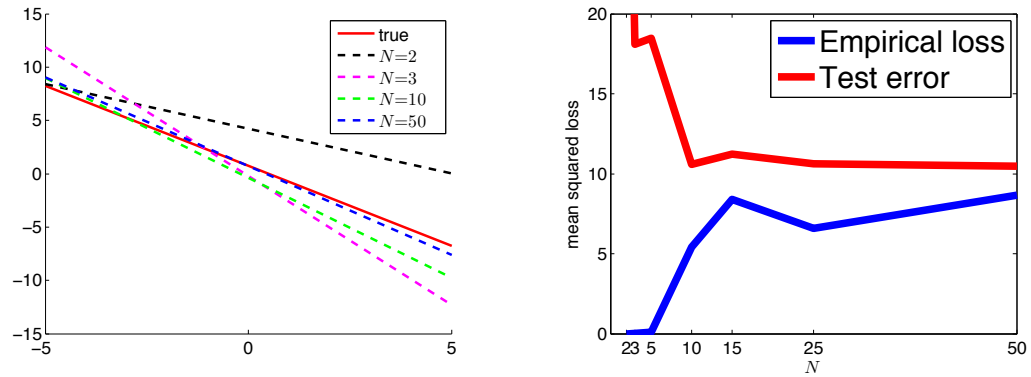
- What happens when we only have a single data point (in 1D)?
 - Ill-posed problem: an infinite number of lines pass through the point and produce “perfect” fit.
- Two points in 1D?
- Two points in 2D?
- This is a general phenomenon: the amount of data needed to obtain a meaningful estimate of a model is related to the number of parameters in the model (its *complexity*).

22



Linear regression - generalization

- Toy experiment: fit a line to varying number of points drawn from the same distribution $p(\mathbf{x}, y)$



- A paradox?
 - The more training data we have, the “worse” is the fit;
 - But at the same time our prediction ability seems to improve.