

# EN.600.475 Machine Learning

## Linear Regression

Raman Arora  
Lecture 4  
February 9, 2017

- Loss and Risk
- Least squares estimation

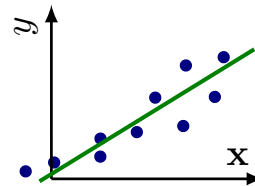
Slides credit: Greg Shakhnarovich <sup>1</sup>

Review

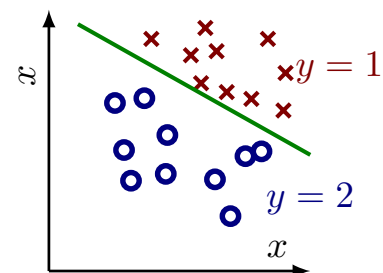
## Review: supervised learning

- Task: build a mapping from input  $\mathcal{X}$  output  $\mathcal{Y}$
- Given a training set  $(\mathbf{x}_i, y_i)$ ,  $i = 1, \dots, N$ , with  $\mathbf{x}_i \in \mathcal{X}$ ,  $y_i \in \mathcal{Y}$ .
- Goal (informally): predict  $y$  accurately for future  $x$ s

regression:  $\mathcal{Y} = \mathbb{R}$   
learn a (continuous) function  $f$



classification:  $\mathcal{Y} = \{1, \dots, C\}$   
learn a separator between classes



## Review: supervised learning pipeline

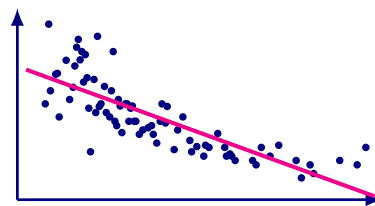
- **Set up (define)** a supervised learning problem
- **Data collection** for training and test set.
- **Representation** choose how data are fed to the model
- **Modeling** Choose a *hypothesis (model) class*
- **Estimation (learning)** Find best hypothesis you can in the chosen class, given the data.
- **Model (class) selection**

3

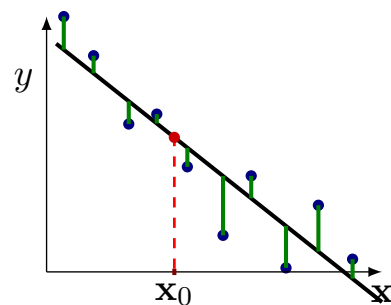


## Review: linear regression

- Two goals in mind:  
Explain the data  
Make predictions



- Model class: linear functions
- Fitting criterion, to guide selection of a function: sum of squared distances from data to the line, along  $y$  axis



9



# Roadmap

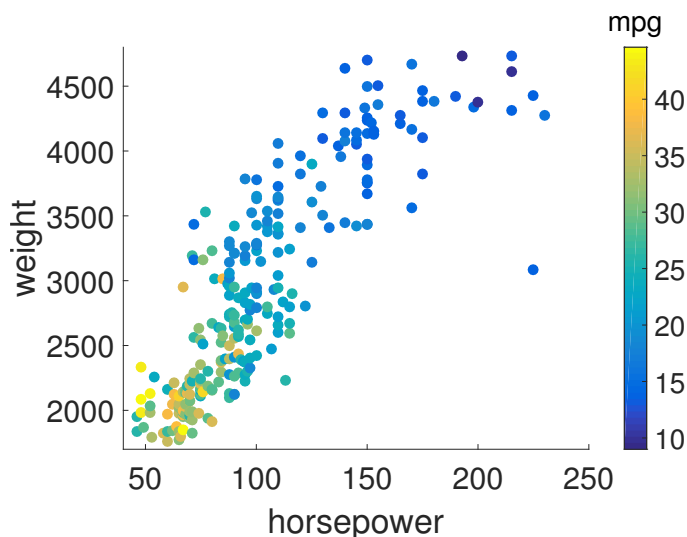
- General form of linear regression and least squares fit
- Loss and risk: definitions and analysis
- Analysis of error in empirical risk minimization

10



## Multiple input variables

- Can consider additional features; e.g.,  $x_1$  horsepower and  $x_2$  vehicle weight.
- We now have mapping from  $\mathbf{x} = (x_1, x_2) \in \mathbb{R}^2$  to  $y$



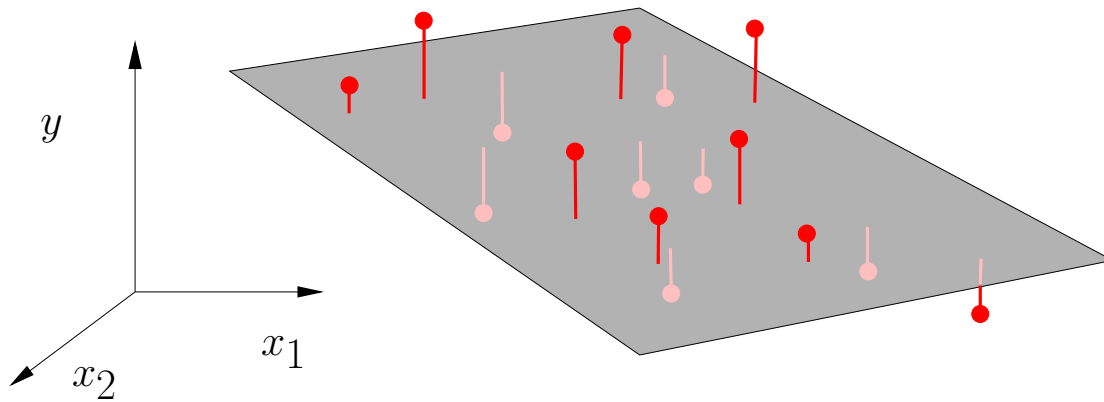
11

colorbar: one possible way to convey multi-dimensional plots



## Fitting a plane to data

- Can use the same criterion: minimize sum of square distances along  $y$ -axis



12



## Linear functions

- General form:

$$\begin{aligned} f(\mathbf{x}; \mathbf{w}) &= w_0 + w_1x_1 + \dots + w_dx_d \\ &= \mathbf{w} \cdot \mathbf{x} \end{aligned}$$

denoting  $x_0 \equiv 1$

- 1D case ( $\mathcal{X} = \mathbb{R}$ ): a line
- $\mathcal{X} = \mathbb{R}^2$ : a plane
- *Hyperplane* in general,  $d$ -D case.

13



## Notation

We will mostly stick to these throughout the course:

- $\mathbf{x}_i$  the  $i$ -th data point in  $\mathcal{X}$  (column vector)  
Often  $\mathcal{X} \equiv \mathbb{R}^d$ , so that  $\mathbf{x}_i = [x_{i1}, \dots, x_{id}]$   
Often assume also  $x_{i0} \equiv 1$
- $y_i$  the label of the  $i$ -th data point;  $y_i \in \mathcal{Y}$
- $\mathbf{x}_0, y_0$  a single test point and its (unknown) label
- $\mathbf{X}$  the  $N \times d$  data matrix where  $i$ -th row is  $\mathbf{x}_i$
- $\mathbf{y}$  the label vector  $\mathbf{y} = [y_1, \dots, y_N]$
- $\mathbf{w} \cdot \mathbf{x}$  inner (dot) product,  $\sum_j w_j x_j$   
sometimes write  $\mathbf{w}^T \mathbf{x}$

14



## Loss function

- Recall: target labels are in  $\mathcal{Y}$  (e.g., regression:  $\mathcal{Y} \equiv \mathbb{R}$ )
- A *loss function*  $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$  maps prediction to cost, given true value:  
 $\ell(\hat{y}, y)$  defines the penalty paid for predicting  $\hat{y}$  when the true value is  $y$ .
- Standard choice for regression: squared loss  $\ell(\hat{y}, y) = (\hat{y} - y)^2$   
is it a good loss function?..
- It is symmetric (sign of mistake doesn't matter); non-negative; gives zero loss for correct prediction
- Vaguely justifiable as “energy” of something
- Penalizes quite harshly for larger mistakes

17



## Empirical risk

- We consider a *parametric* function  $f(\mathbf{x}; \mathbf{w})$   
E.g., linear function:  $f(\mathbf{x}; \mathbf{w}) = w_0 + \sum_{j=1}^d w_j x_{ij} = \mathbf{w} \cdot \mathbf{x}_i$
- The *empirical* risk of function  $y = f(\mathbf{x}; \mathbf{w})$  on a set  $\mathbf{X}$ :

$$\hat{R}_n(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \ell(f(\mathbf{x}_i; \mathbf{w}), y_i)$$

- LSQ minimizes the empirical risk when  $\ell$  is squared loss.
- We care about accuracy of *predicting* labels for new examples.  
Why/when does empirical risk minimization help us achieve that?

18



## Risk: empirical and expected

- Fundamental assumption: example  $\mathbf{x}$ /label  $y$  are drawn from a joint probability distribution  $p(\mathbf{x}, y)$ .
- Data are i.i.d.: same (unknown!) distribution for all pairs  $(\mathbf{x}, y)$  in both training and test data.
- We can measure the empirical risk on training set

$$\hat{R}_n(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \ell(f(\mathbf{x}_i; \mathbf{w}), y_i)$$

- The ultimate goal is to minimize the *expected loss*, also known as *risk*:

$$R(\mathbf{w}) = E_{(\mathbf{x}_0, y_0) \sim p(\mathbf{x}, y)} [\ell(f(\mathbf{x}_0; \mathbf{w}), y_0)]$$

19



## Risk: empirical and expected

- Empirical risk:

$$\hat{R}_n(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \ell(f(\mathbf{x}_i, \mathbf{w}), y_i)$$

- Risk:

$$R(\mathbf{w}) = E_{(\mathbf{x}_0, y_0) \sim p(\mathbf{x}, y)} [\ell(f(\mathbf{x}_0, \mathbf{w}), y_0)]$$

- Empirical risk minimization (ERM) approach: to the extent that the training set is a representative of the underlying distribution  $p(\mathbf{x}, y)$ , the empirical loss serves as a proxy for the risk (expected loss).
- Technically: estimate  $p(\mathbf{x}, y)$  by the *empirical distribution* of data.

20



## Learning via empirical risk minimization

Two steps:

- Select a restricted class  $\mathcal{H}$  of *hypotheses*  $f : \mathcal{X} \rightarrow \mathcal{Y}$   
Here: linear functions parametrized by  $\mathbf{w}$ :  $f(\mathbf{x}; \mathbf{w}) = \mathbf{w} \cdot \mathbf{x}$
- Select a hypothesis  $f^* \in \mathcal{H}$  based on training set  
Here: minimize empirical squared loss, i.e., select  $f(\mathbf{x}; \mathbf{w}^*)$  where

$$\hat{\mathbf{w}} = \operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

- How do we find  $\hat{\mathbf{w}} = [\hat{w}_0, \hat{w}_1, \dots, \hat{w}_d]$  ?

21



## Least squares: estimation

- We need to minimize  $\hat{R}_n(\mathbf{w})$  w.r.t.  $\mathbf{w}$

$$\hat{R}_n(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

- Necessary condition to minimize  $\hat{R}_n(\mathbf{w})$ :

$$\frac{\partial \hat{R}_n(\mathbf{w})}{\partial \mathbf{w}} = \mathbf{0},$$

i.e., derivatives w.r.t.  $w_0, w_1, \dots, w_d$  must all be zero.

22



## Matrix derivatives

- Scalar valued function of one variable

$$f : \mathbb{R} \rightarrow \mathbb{R} \quad \text{derivative: } \frac{df}{dx}$$

- Scalar valued function of multiple scalar variables

$$f : \underbrace{\mathbb{R} \times \dots \times \mathbb{R}}_{d \text{ times}} \rightarrow \mathbb{R} \quad \text{gradient: } \nabla f = \left[ \frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_d} \right]$$

- If we collect multiple variables in a vector:  $\mathbf{x} \in \mathbb{R}^d$ :

$$\nabla f = \frac{\partial f}{\partial \mathbf{x}}$$

derivative of  $f$  w.r.t.  $\mathbf{x}$  has the same dimension as  $\mathbf{x}$

23





## Least squares: estimation

$$\hat{R}_n(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

$$\frac{\partial \hat{R}_n(\mathbf{w})}{\partial w_0} = \frac{1}{N} \sum_{i=1}^N \frac{\partial}{\partial w_0} (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2 = 0$$

$$\Rightarrow \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) = 0$$

- $y_i - \mathbf{w} \cdot \mathbf{x}_i$  is the *prediction error* on the  $i$ -th example.
- $\Rightarrow$  Necessary condition for optimal  $\mathbf{w}$  is that the errors have zero mean. (why does it make sense?)

26



## Least squares: estimation

$$\hat{R}_n(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i)^2$$

$$\frac{\partial \hat{R}_n(\mathbf{w})}{\partial w_j} = - \frac{2}{N} \sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) x_{ij} = 0.$$

- Second necessary condition: errors are *uncorrelated* with the data!  
(And with *any linear function* of the data)
- $d + 1$  linear equations in  $d + 1$  unknowns  $w_0, w_1, \dots, w_d$

$$\sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) x_{ij} = 0 \quad \forall j = 1, \dots, d, \quad (1)$$

$$\sum_{i=1}^N (y_i - \mathbf{w} \cdot \mathbf{x}_i) = 0 \quad (2)$$

28



## Least squares in matrix form

$$\mathbf{X} = \begin{bmatrix} 1 & x_{11} & \cdots & x_{1d} \\ \vdots & & & \\ 1 & x_{N1} & \cdots & x_{Nd} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} w_0 \\ \vdots \\ w_d \end{bmatrix}$$

- Predictions:  $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$ , errors:  $\mathbf{y} - \mathbf{X}\mathbf{w}$ , empirical loss:

$$\begin{aligned} \hat{R}_n(\mathbf{w}) &= \frac{1}{N} (\mathbf{y} - \mathbf{X}\mathbf{w}) \cdot (\mathbf{y} - \mathbf{X}\mathbf{w}) = \frac{1}{N} (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w}) \\ &= \frac{1}{N} (\mathbf{y}^T - \mathbf{w}^T \mathbf{X}^T) (\mathbf{y} - \mathbf{X}\mathbf{w}) \end{aligned}$$

Using  $(AB)^T = B^T A^T$ ,  $(A + B)^T = A^T + B^T$ ,  $(A^T)^T = A$ .

30



## Derivative of loss

$$\hat{R}_n(\mathbf{w}) = \frac{1}{N} (\mathbf{y}^T - \mathbf{w}^T \mathbf{X}^T) (\mathbf{y} - \mathbf{X}\mathbf{w}).$$

$$\frac{\partial \mathbf{a}^T \mathbf{b}}{\partial \mathbf{a}} = \frac{\partial \mathbf{b}^T \mathbf{a}}{\partial \mathbf{a}} = \mathbf{b}, \quad \frac{\partial \mathbf{a}^T \mathbf{B} \mathbf{a}}{\partial \mathbf{a}} = 2\mathbf{B}\mathbf{a}$$

$$\begin{aligned} \frac{\partial \hat{R}_n(\mathbf{w})}{\partial \mathbf{w}} &= \frac{1}{N} \frac{\partial}{\partial \mathbf{w}} [\mathbf{y}^T \mathbf{y} - \mathbf{w}^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \mathbf{w} + \mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}] \\ &= \frac{1}{N} [\mathbf{0} - \mathbf{X}^T \mathbf{y} - (\mathbf{y}^T \mathbf{X})^T + 2\mathbf{X}^T \mathbf{X} \mathbf{w}] \\ &= -\frac{2}{N} (\mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X} \mathbf{w}) \end{aligned}$$

34



## Least squares solution

$$\frac{\partial \hat{R}_n(\mathbf{w})}{\partial \mathbf{w}} = -\frac{2}{N}(\mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X} \mathbf{w}) = 0$$

$$\mathbf{X}^T \mathbf{y} = \mathbf{X}^T \mathbf{X} \mathbf{w} \Rightarrow \hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- $\mathbf{X}^\dagger \triangleq (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$  is called the *Moore-Penrose pseudoinverse* of  $\mathbf{X}$ .
- Linear regression in Python:  

```
X[:,0]=1; X[:,1::]=x # assumes X is right size
w=np.dot(np.linalg.pinv(X),y)
```
- Prediction:  $\hat{y} = \hat{\mathbf{w}} \cdot \mathbf{x}_0$

$$\hat{y} = \hat{\mathbf{w}} \cdot \mathbf{x}_0 = \mathbf{y}^T \mathbf{X}^\dagger \mathbf{x}_0$$

Note: we have  $d + 1$  numbers in  $\hat{\mathbf{w}}$  capture what the training data  $\mathbf{X}$ ,  $\mathbf{y}$  tell us about  $\mathcal{X} \rightarrow \mathcal{Y}$  *under our model class*

39



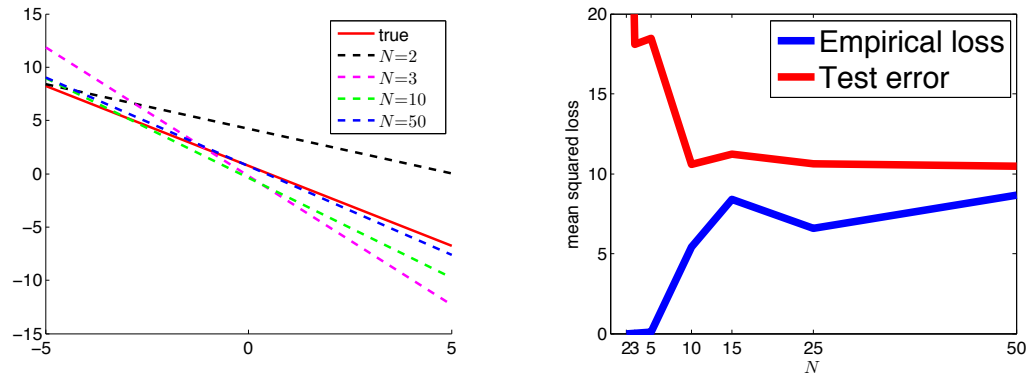
## Data set size and regression

- What happens when we only have a single data point (in 1D)?
  - Ill-posed problem: an infinite number of lines pass through the point and produce “perfect” fit.
- Two points in 1D?
- Two points in 2D?
- This is a general phenomenon: the amount of data needed to obtain a meaningful estimate of a model is related to the number of parameters in the model (its *complexity*).



# Linear regression - generalization

- Toy experiment: fit a line to varying number of points drawn from the same distribution  $p(\mathbf{x}, y)$



- A paradox?
  - The more training data we have, the “worse” is the fit;
  - But at the same time our prediction ability seems to improve.